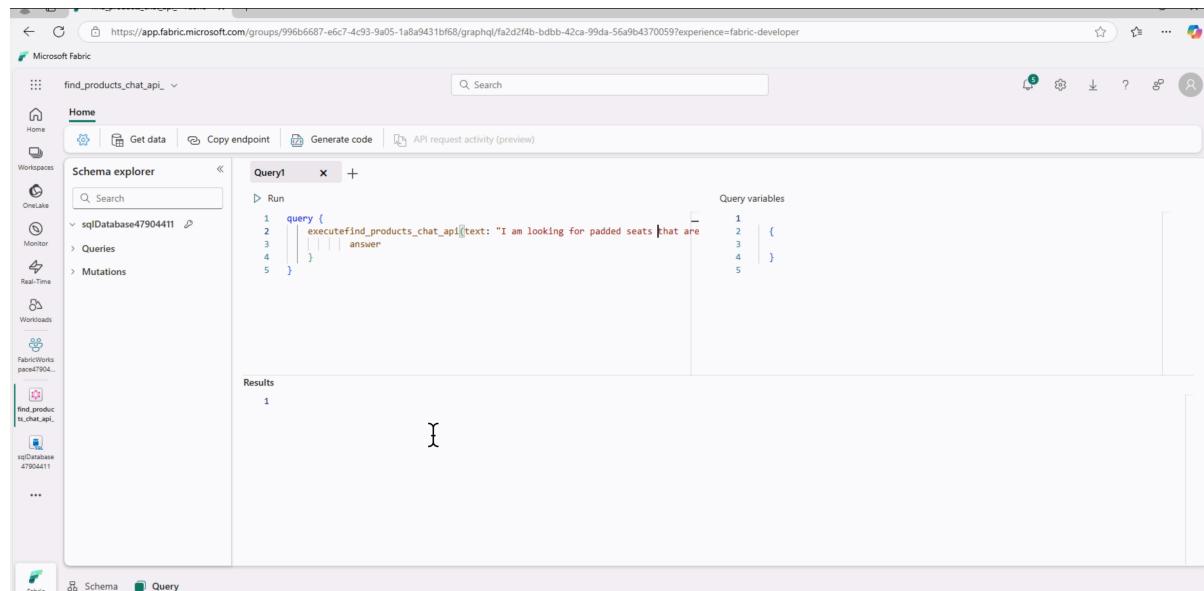


# Build an AI App with SQL DB in Fabric

Today's lab will guide you through creating a set of GraphQL RAG application APIs that use relational data, Azure OpenAI, and SQL DB in Fabric.



The screenshot shows the Microsoft Fabric Query interface. The URL is https://app.fabric.microsoft.com/groups/996b6687-e6c7-4c93-9a05-1a8a9431bf68/graphql/fa2d2f4b-bdbb-42ca-99da-56a9b4370059?experience=fabric-developer. The left sidebar shows 'Workspaces' with 'OneLake' selected, and 'Queries' under 'sqlDatabase47904411'. The main area has a 'Query1' tab open with the following GraphQL code:

```
query {
  executeFind_products_chat_api(text: "I am looking for padded seats. What are they like?")
}
```

The results pane shows a single row of data:

text	answer
I am looking for padded seats. What are they like?	[REDACTED]

## Key Learning Objectives

- Query and work with data in SQL DB in Microsoft Fabric
- Build AI application design patterns/endpoints
- Use AI with relational data
- Understand how similarity searching benefits queries

# Lab Overview

1. The Azure Fabric Portal and creating a SQL database in Fabric
2. Working with the SQL Database in Microsoft Fabric
3. REST in the database
4. Creating embeddings for relational data with Azure OpenAI
5. Create a set of GraphQL endpoints for AI Applications
6. Create a PowerBI report with Copilot
7. DevOps in Fabric (demo only)



# **The Azure Fabric Portal and creating a SQL database in Fabric**



# SQL Database in Fabric

Autonomous database built  
on the foundation of Azure SQL Database



## Simple

Integrated with Visual Studio and GitHub  
Developer optimized with GraphQL  
and CI/CD support  
Unify OLAP + OLTP



## Autonomous

Deploy in seconds  
Auto-scaled compute and storage  
Auto-managed HADR  
Auto-optimized performance



## Integrated

Integration with Azure AI  
Integrated with Fabric services  
Built-in vector and RAG support  
AI-assisted



## Secure by default



## Powered by SQL



## OneLake integrated

# Beyond RDBMS



<https://aka.ms/sqldev>

{ } JSON and RegEx

Spatial

Graph

Column store

Memory-optimized

Ledger

Query Intelligence

{ } REST API integrated

Data API Builder, GraphQL

Azure Functions

Hyperscale, Serverless, Pools

Containers, DevOps, GitHub

AI

Fabric

# Chapter 1 Objectives

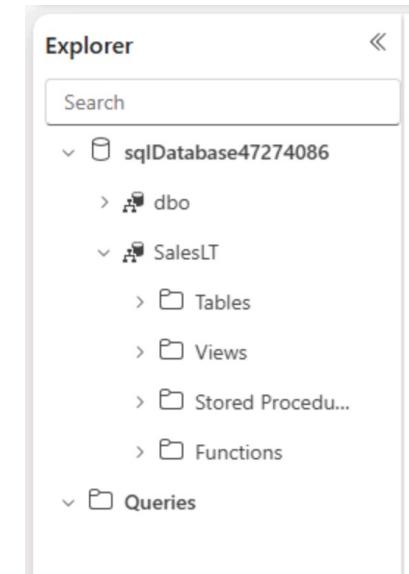
- Log into Microsoft Fabric
- Create a workspace
- Create a SQL DB
- Load the database with the Adventure Works sample data

# **Working with the SQL Database in Microsoft Fabric**



# SQL query editor in the Fabric portal

- View database schemas and objects via the database explorer
- Write and execute T-SQL queries
- Work with multiple editors/queries via tabs
- Use templates to jumpstart development
- Performance Monitor
- Save a query as a view
- Download results as Excel/JSON/CSV files
- Create GraphQL endpoints



# Copilot for SQL database in Microsoft Fabric

- Natural Language to SQL
- Document-based Q&A
- Auto Code Generation
- Fix Query Errors
- Explain Query



# Chapter 2 Objectives

- Use the database explorer in Microsoft Fabric
- Run SQL statements in the SQL query editor
- Get answers with Copilot for SQL database in Microsoft Fabric

**Note: Syntax highlighting may sometimes indicate errors, please ignore this.**

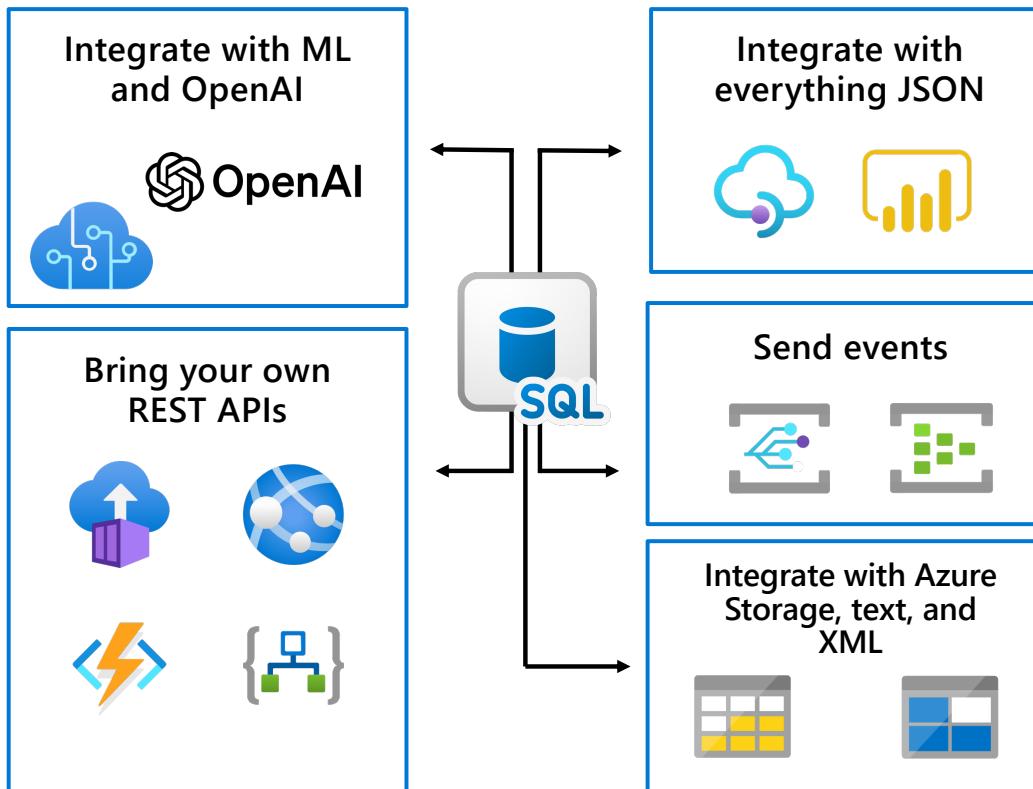
# **Getting started with REST in the database**







# SQL Database and External REST endpoints



[sp\\_invoke\\_external\\_rest\\_endpoint](#)

EXECUTE ANY EXTERNAL ENDPOINT

# External REST Endpoint Invocation (EREI)

## sp\_invoke\_external\_rest\_endpoint

```
declare @url nvarchar(4000) = N'https://chat.openai.azure.com/';
declare @payload nvarchar(max) = N'{"messages": [{"role": "system", "content": "XYZ"}, {"role": "system", "content": "ABC"}]}';
declare @ret int, @response nvarchar(max);

exec @ret = sp_invoke_external_rest_endpoint
    @url = @url,
    @method = 'POST',
    @payload = @payload,
    @credential = [https://chat.openai.azure.com/],
    @timeout = 230,
    @response = @response output;

select @response;
```

# Chapter 3 Objectives

- Create database scoped credentials
- Test the connection to Azure OpenAI from in the database
  - Chat completion
  - Content safety
  - AI Translator



**Creating embeddings  
for relational data with  
Azure OpenAI**

# Azure OpenAI Service

Azure OpenAI Service provides REST API access to OpenAI's powerful language models.

Models	Description
o1-preview and o1-mini	Designed to tackle reasoning and problem-solving tasks with increased focus and capability (In Preview).
GPT-4o & GPT-4o mini & GPT-4 Turbo	The latest most capable Azure OpenAI models with multimodal versions, which can accept both text and images as input.
GPT-4o audio	A GPT-4o model that supports low-latency, "speech in, speech out" conversational interactions.
GPT-4	A set of models that improve on GPT-3.5 and can understand and generate natural language and code.
Embeddings	A set of models that can convert text into numerical vector form to facilitate text similarity.
DALL-E	A series of models that can generate original images from natural language.

# REST and Inferencing Endpoints

## Embeddings

```
{ "input": [ "this is a test" ] }
```



```
{ "body": { "data":  
[ { "index": 0, "embedding": [ 0.13, 1.43, ... ] } ] } }
```

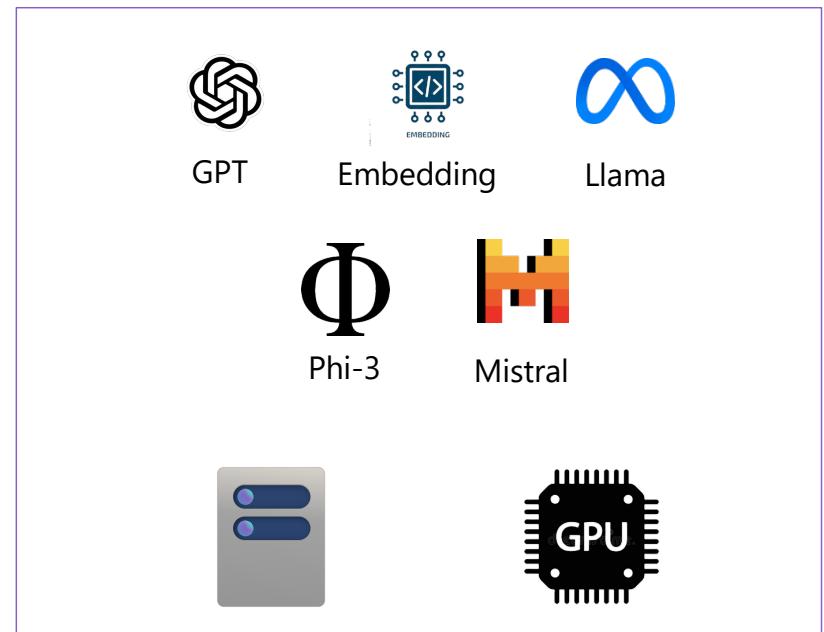


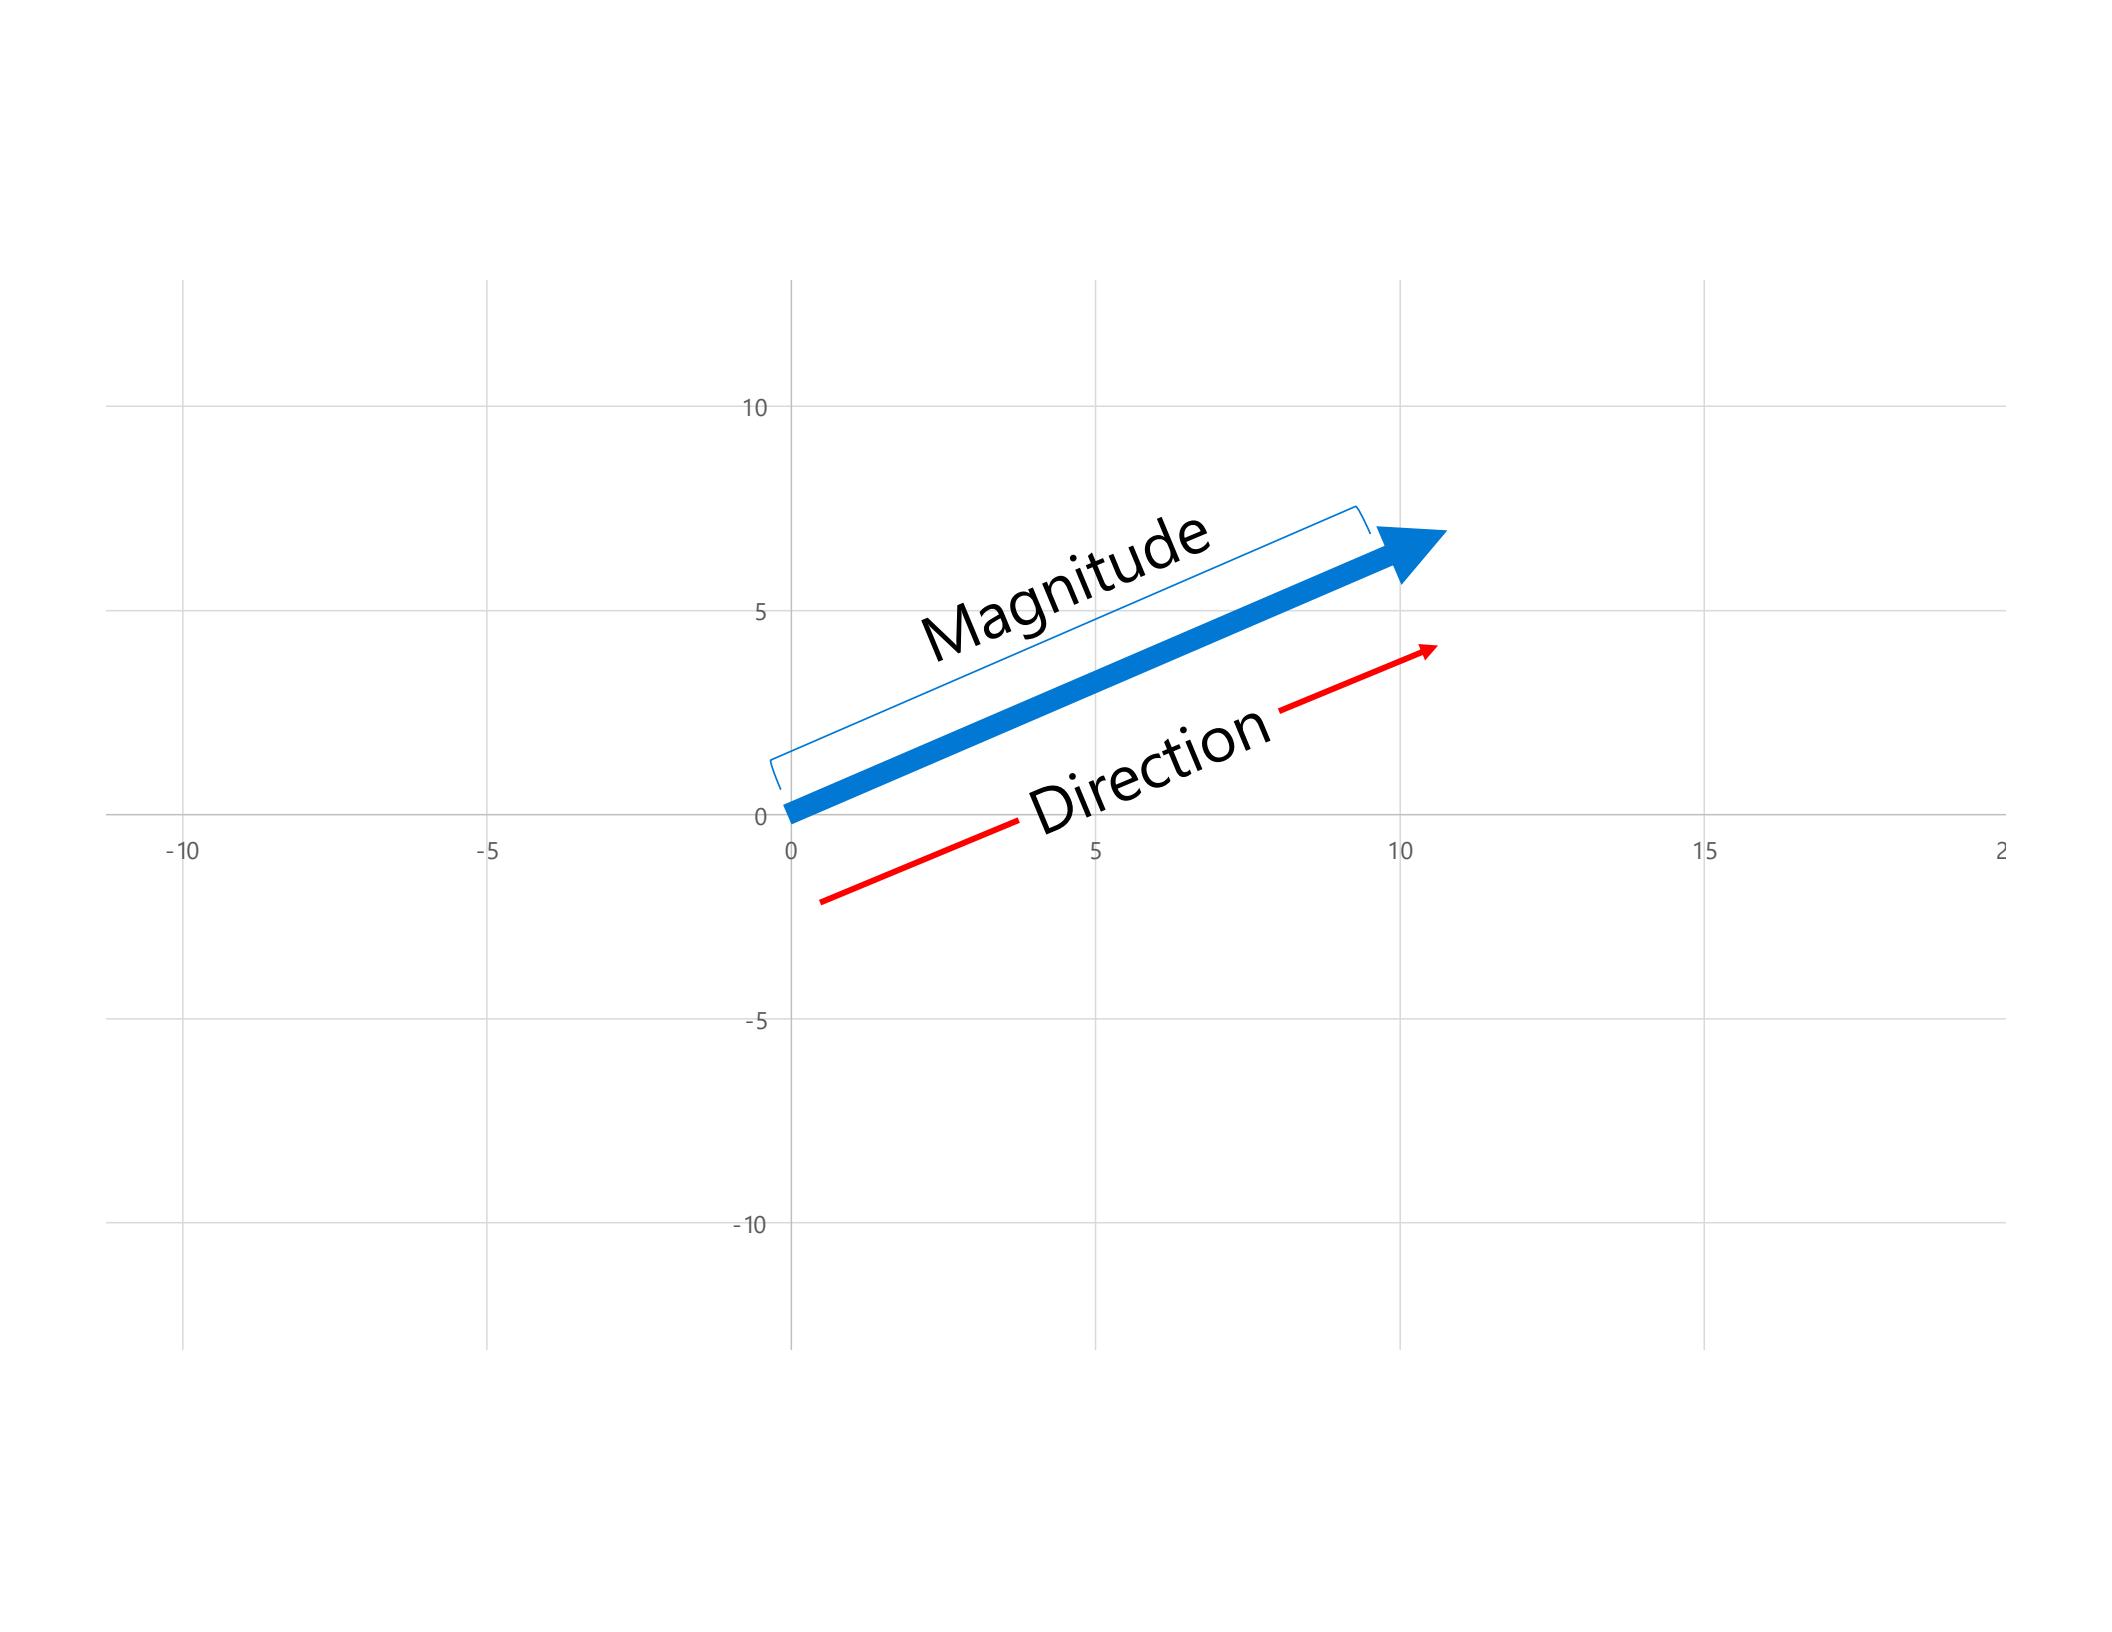
## Chat Completions

```
[  
  {"role": "system", "content": "You are a helpful assistant."},  
  {"role": "user", "content": "Who founded Microsoft?"}  
]
```



```
[  
  {"role": "assistant", "content": "Microsoft was founded by Bill Gates and Paul Allen."}  
]
```





Magnitude

Direction

How many stars did each user give these 3 movies?

User 1: [5, 3, 3]

User 2: [1, 3, 4]

# Vectors and Embeddings

## What is vector embedding?

A **Vector Embedding** is a technique used in machine learning to transform data (text, images, etc.) into numerical representations (vectors) that capture the essential characteristics and relationships of the data.

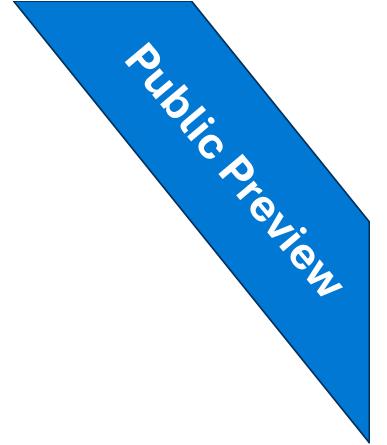
Each number is referred to as a dimension.

dog: [0.9, 0.3, 0.2, ...]

puppy: [0.88, 0.33, 0.21, ...]

Image of a dog: [0.9, 0.3, 0.2, ...]

"I took the dog for a walk": [1.5, -0.8, 2.1, ...]

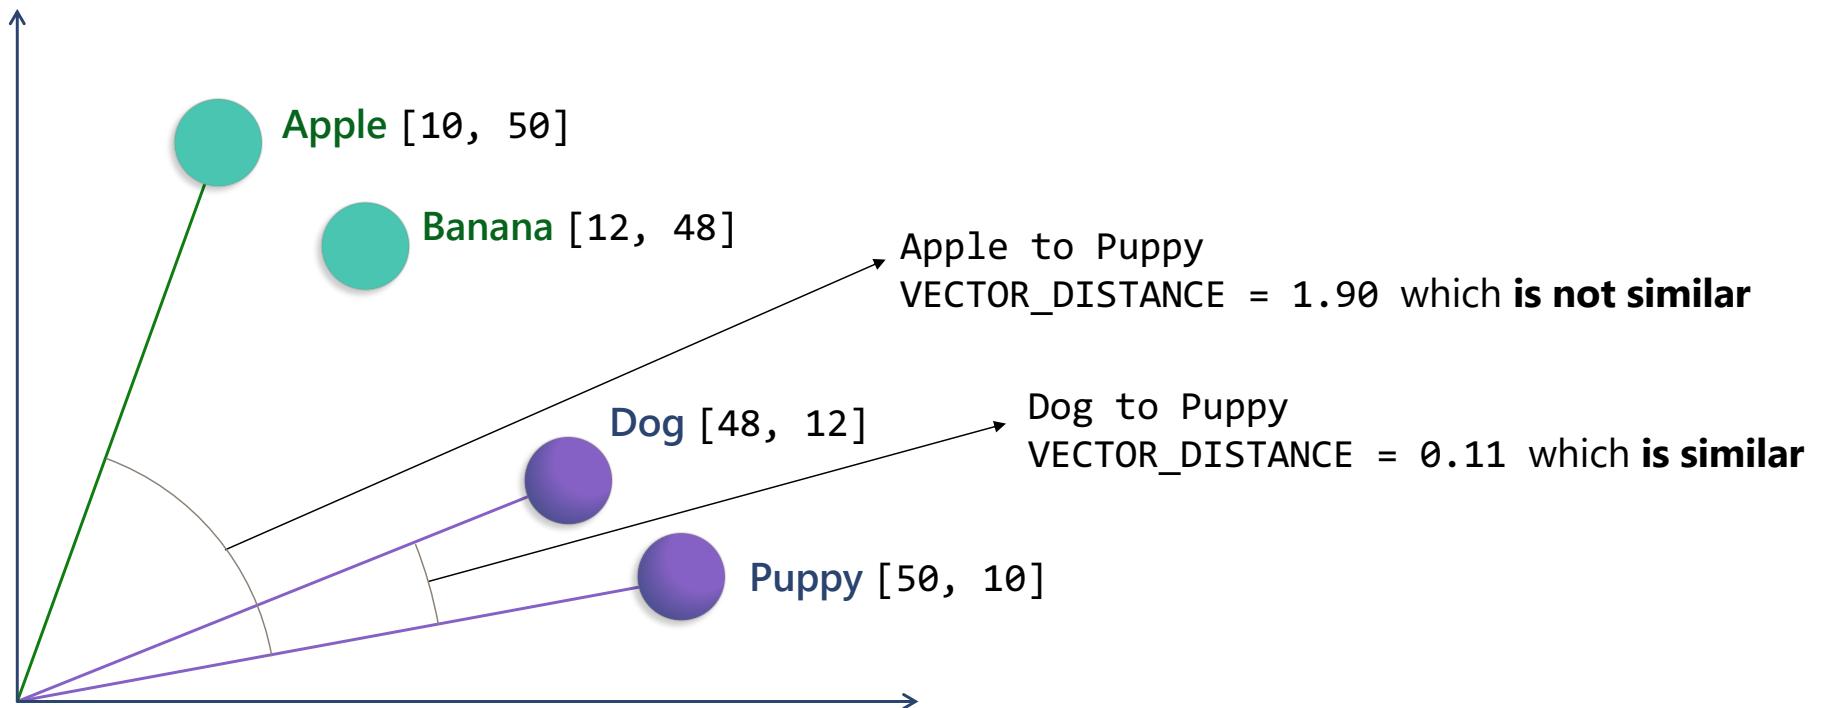


Public Preview

## New in SQL Database in Fabric

- Vector Datatype
  - col\_name VECTOR(`int`)  
*(int being the dimension count)*
- Vector Functions
  - VECTOR\_DISTANCE
  - VECTOR\_NORM
  - VECTOR\_NORMALIZE

# Similarity Searching using VECTOR\_DISTANCE (cosine)



# Similarity Searches vs Text Only Searches

**“I’m looking to buy a comfy red sofa”**

## Text Search Results

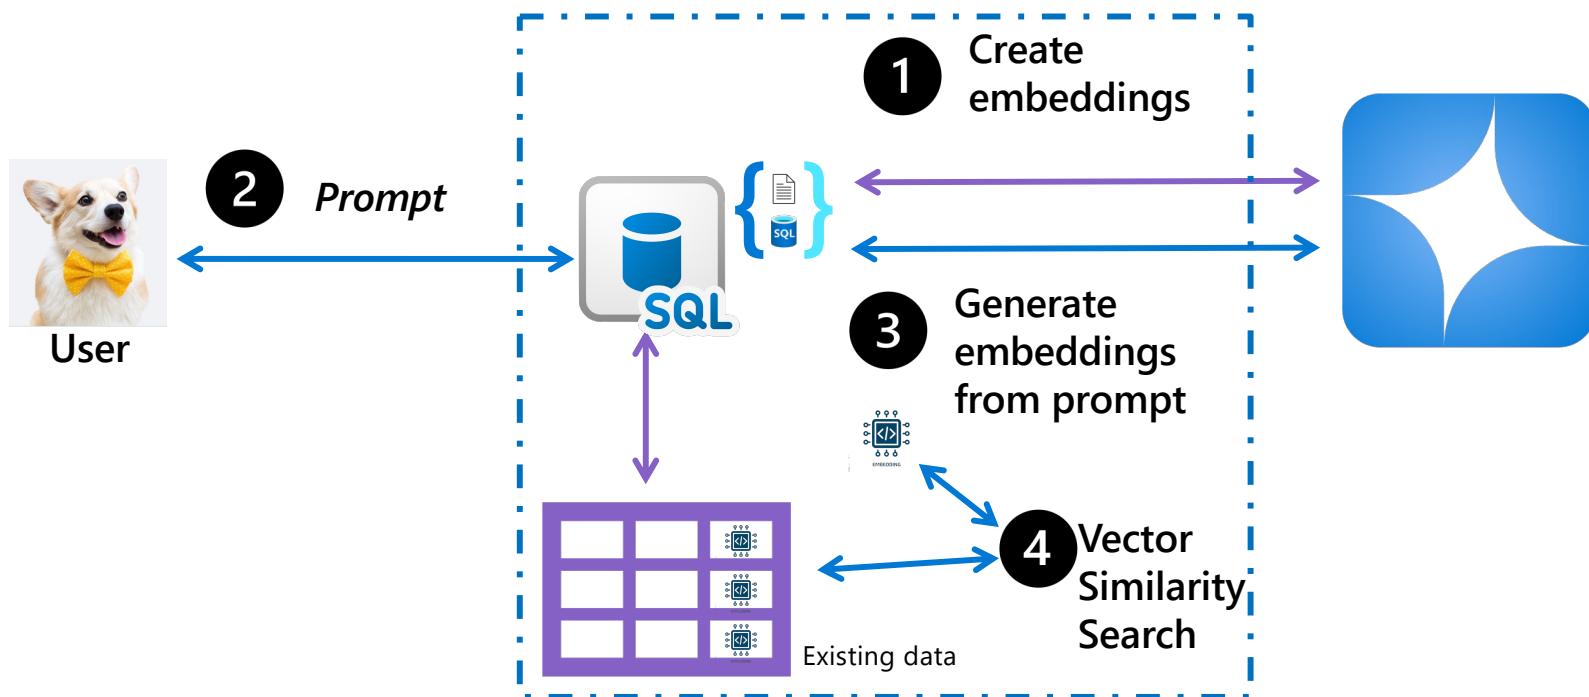
Product Name
Red Sofa
Comfy Sofa - Red

## Similarity Search Results

Product Name
Red Sofa
Comfy Sofa - Red
Soft maroon couch
Sleepy time crimson daybed

# Use T-SQL for Vector and Hybrid Search

Connected to AI



## How do I handle changes to the data?

- Database Change Tracking/Change Data Capture
- SQL Trigger Binding for Azure Functions
- Batch Jobs
- Event Architectures

# Chapter 4 Objectives

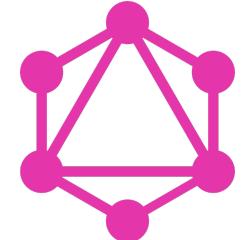
- Add columns to the Products table for embeddings
- Create a stored procedure to embed table data
- Create embeddings for the Products table
- Use the new VECTOR\_DISTANCE function for similarity searches

# Create GraphQL endpoints for AI Applications



# What is GraphQL

Ask for what you need, get exactly that



Clients specify exactly what data they need, and the server returns only that data.

```
{  
  user(id: "1") {  
    name  
    email  
  }  
}
```



```
{  
  "data": {  
    "user": {  
      "name": "Frank",  
      "email": "frank@frank.com"  
    }  
  }  
}
```

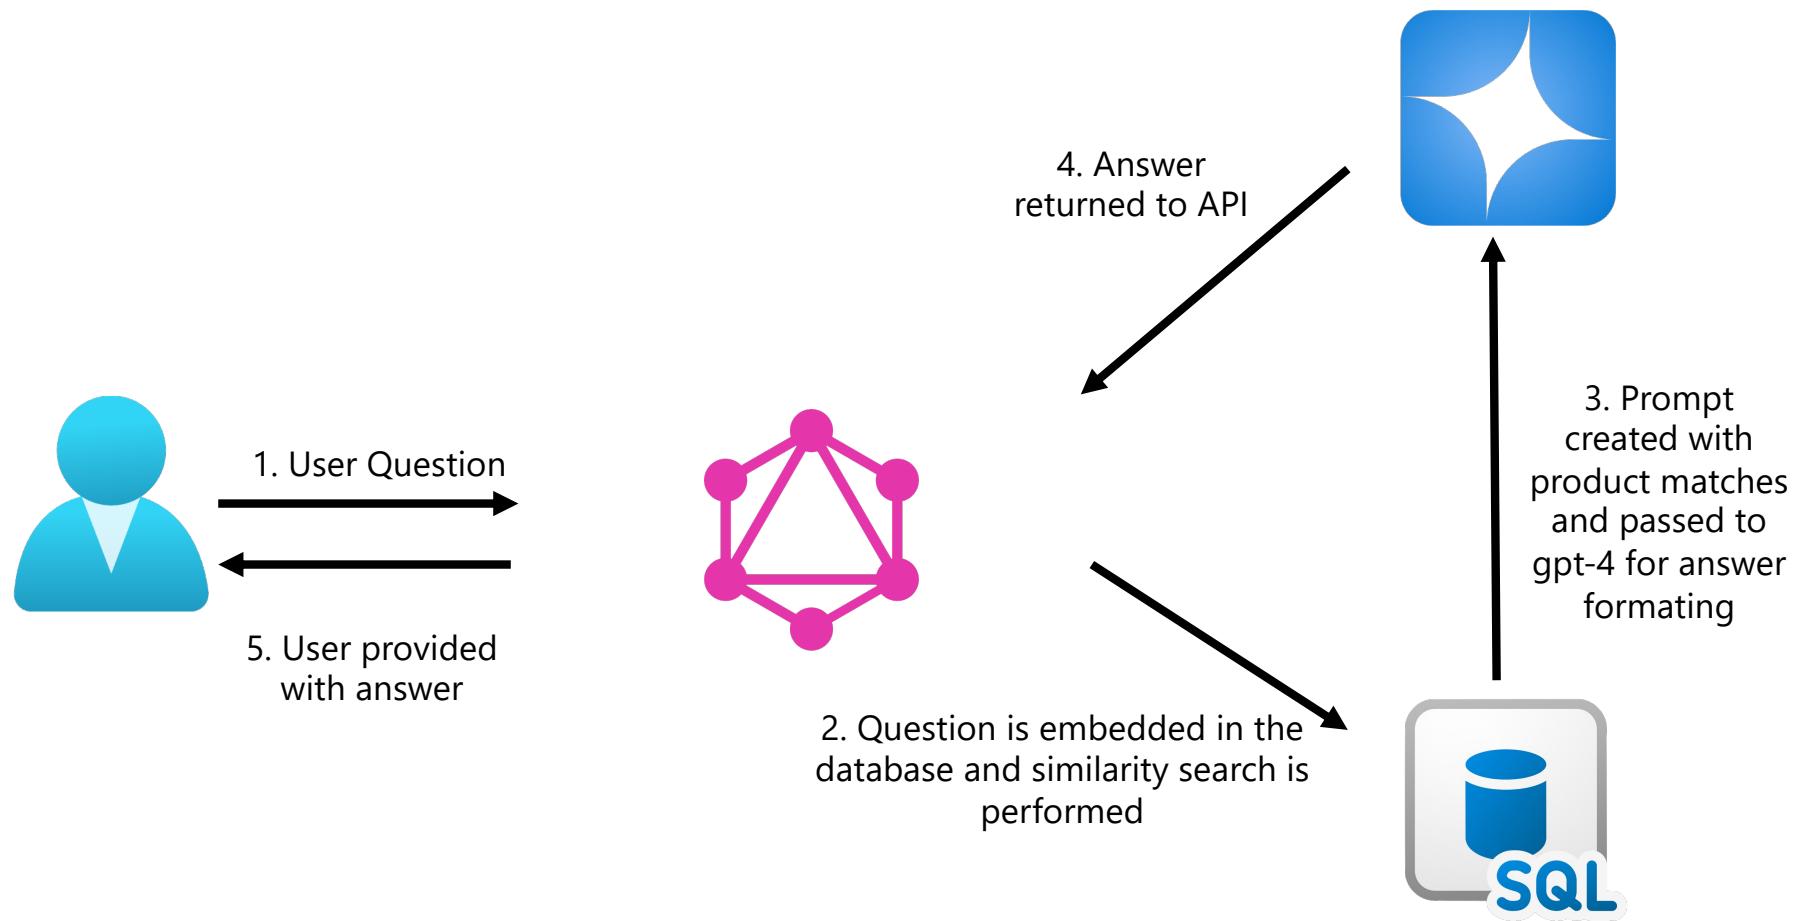
# **Retrieval-Augmented Generation and AI**

**Retrieval-Augmented Generation** (RAG) is a significant advancement in AI because it enhances the capabilities of large language models (LLMs) by integrating external knowledge sources such as relational databases.

## **Key points:**

- Improved Accuracy and Relevance
- Cost-Effective
- Greater Control for Developers

# A GraphQL API using RAG and SQL



# Chapter 5 Objectives

- Create a stored procedure for a RAG application
- Create a GraphQL API in Fabric and test it
- Add chat completion to our RAG stored procedure
- Create a GraphQL API in Fabric to expose the new chat function

# **Creating a PowerBI Report from a SQL Database in Microsoft Fabric with Copilot**



# Copilot for Power BI

Copilot for Power BI can help you:

- Get started on a new report by suggesting topics based on your data.
- Use the measures and columns in your semantic model to help you explore your data
- Builds a visual for you if the answer can't be found in the report

# Chapter 6 Objectives

- Add the SQL Database to a semantic model
- Ask Copilot for report suggestions
- Explore the PowerBI dashboard Copilot created from a SQL database

# Azure DevOps and the SQL Database in Microsoft Fabric



# Microsoft Fabric Git integration

## Azure DevOps or GitHub

- Backup and version work
- Revert to previous stages as needed
- Collaborate with others or work alone using Git branches
- Apply the capabilities of familiar source control tools to manage Fabric items

### Supported Services:

Data pipelines (preview)	Lakehouse (preview)	Reports Semantic models (preview)
Dataflows gen2 (preview)	Mirrored database (preview)	Spark Job Definitions (preview)
Eventhouse and KQL database (preview)	Notebooks	Spark environment (preview)
EventStream (preview)	Paginated reports (preview) Reflex (preview)	<b>SQL database (preview)</b> Warehouses (preview)